

# Exploratory Data Analysis Problem Set 3

Robert Chen

December 13, 2015

## Summary


This is a summary of the code to complete Problem Set 3 of UD561 for the Exploratory Data Analysis Unit (Part 1) of the Sliderule "Fundamentals of Data Science" course.

## Problem 1: Diamonds -- Price per Carat vs Color

The first part of the assignment is to explore price per carat vs color in the "diamonds" dataset using boxplots ("diamonds" comes with the library "ggplot2"). color ranges from D to J, with D being the best grade of color.

The code follows:

```
library(ggplot2)
diamonds$ppc <- NA
diamonds$ppc <- diamonds$price / diamonds$carat
qplot(x = color, y = ppc, data = diamonds, geom="boxplot",
      xlab = "Color",
      ylab = "Price per Carat")
```



The first line loads the ggplot2 library (including the "diamonds" dataset). The next 2 lines create a price per carat variable. The third line creates the box plots.

This can also be sanity-checked using the "summary" command using the ppc variable across color:

```
by(diamonds$ppc, diamonds$color, summary)

## diamonds$color: D
##   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##   1128   2455   3411   3953   4749   17830
## -----
## diamonds$color: E
##   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##   1078   2430   3254   3805   4508   14610
## -----
## diamonds$color: F
##   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##   1168   2587   3494   4135   4947   13860
## -----
## diamonds$color: G
##   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##   1139   2538   3490   4163   5500   12460
## -----
## diamonds$color: H
##   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##   1051   2397   3819   4008   5127   10190
```


```
## -----
## diamonds$color: I
##   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##   1152   2345   3780   3996   5197   9398
## -----
## diamonds$color: J
##   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##   1081   2563   3780   3826   4928   8647
```

From the data, one can see that all colors roughly have the same median and middle quantiles, but the maximum value increases as the color grade gets better.

## Problem 2: Diamonds -- Distribution of Carats


The second part is to explore weight (carats) using frequency polygons. First we try with a binwidth of 10:

```
qplot(x = carat, data = diamonds, binwidth = 10, geom="freqpoly")
```




This data isn't very helpful. Let's try lowering the binwidth to 1:

```
qplot(x = carat, data = diamonds, binwidth = 1, geom="freqpoly")
```



That's a little better, but it still seems like information is lost. Let's try to go by tenths:


```
qplot(x = carat, data = diamonds, binwidth = .1, geom="freqpoly")
```



This is better. To better answer the question of which carats have quantities greater than 2000 we cause the x-axis values to increment by 0.1 and focus on the area between 0 and 2.5. We also decrease the binwidth to .01 to be able to check the value 1.01 carats:

```
qplot(x = carat, data = diamonds, binwidth = .01, geom="freqpoly") +  
  scale_x_continuous(lim = c(0, 2.5), breaks = seq(0, 2.5, .1))
```


## Warning: Removed 2 rows containing missing values (geom\_path).



From this, one can see that values greater than 2000 start after .25 and stop after 1.05. We then zoom in on that portion and increase the resolution of the x-axis to 0.05:

```
qplot(x = carat, data = diamonds, binwidth = .01, geom="freqpoly") +  
  scale_x_continuous(lim = c(0.25, 1.05), breaks = seq(0.25, 1.05, .05))
```

## Warning: Removed 2 rows containing missing values (geom\_path).



Roughly we see that 0.3 and 1.01 have values above 2000.

This can be confirmed numerically using the table command:

```
table(diamonds$carat)
```

```
##
##  0.2 0.21 0.22 0.23 0.24 0.25 0.26 0.27 0.28 0.29  0.3 0.31 0.32 0.33 0.34
##   12   9   5 293 254 212 253 233 198 130 2604 2249 1840 1189 910
## 0.35 0.36 0.37 0.38 0.39  0.4 0.41 0.42 0.43 0.44 0.45 0.46 0.47 0.48 0.49
##  667 572 394 670 398 1299 1382 706 488 212 110 178 99 63 45
##  0.5 0.51 0.52 0.53 0.54 0.55 0.56 0.57 0.58 0.59  0.6 0.61 0.62 0.63 0.64
## 1258 1127 817 709 625 496 492 430 310 282 228 204 135 102 80
## 0.65 0.66 0.67 0.68 0.69  0.7 0.71 0.72 0.73 0.74 0.75 0.76 0.77 0.78 0.79
##   65  48  48  25  26 1981 1294 764 492 322 249 251 251 187 155
##  0.8 0.81 0.82 0.83 0.84 0.85 0.86 0.87 0.88 0.89  0.9 0.91 0.92 0.93 0.94
##  284 200 140 131 64 62 34 31 23 21 1485 570 226 142 59
## 0.95 0.96 0.97 0.98 0.99  1 1.01 1.02 1.03 1.04 1.05 1.06 1.07 1.08 1.09
##   65 103 59 31 23 1558 2242 883 523 475 361 373 342 246 287
##  1.1 1.11 1.12 1.13 1.14 1.15 1.16 1.17 1.18 1.19  1.2 1.21 1.22 1.23 1.24
##  278 308 251 246 207 149 172 110 123 126 645 473 300 279 236
## 1.25 1.26 1.27 1.28 1.29  1.3 1.31 1.32 1.33 1.34 1.35 1.36 1.37 1.38 1.39
##  187 146 134 106 101 122 133 89 87 68 77 50 46 26 36
##  1.4 1.41 1.42 1.43 1.44 1.45 1.46 1.47 1.48 1.49  1.5 1.51 1.52 1.53 1.54
##   50  40  25  19  18  15  18  21  7  11 793 807 381 220 174
## 1.55 1.56 1.57 1.58 1.59  1.6 1.61 1.62 1.63 1.64 1.65 1.66 1.67 1.68 1.69
##  124 109 106 89 89 95 64 61 50 43 32 30 25 19 24
```

```
## 1.7 1.71 1.72 1.73 1.74 1.75 1.76 1.77 1.78 1.79 1.8 1.81 1.82 1.83 1.84
## 215 119 57 52 40 50 28 17 12 15 21 9 13 18 4
## 1.85 1.86 1.87 1.88 1.89 1.9 1.91 1.92 1.93 1.94 1.95 1.96 1.97 1.98 1.99
## 3 9 7 4 4 7 12 2 6 3 3 4 4 5 3
## 2 2.01 2.02 2.03 2.04 2.05 2.06 2.07 2.08 2.09 2.1 2.11 2.12 2.13 2.14
## 265 440 177 122 86 67 60 50 41 45 52 43 25 21 48
## 2.15 2.16 2.17 2.18 2.19 2.2 2.21 2.22 2.23 2.24 2.25 2.26 2.27 2.28 2.29
## 22 25 18 31 22 32 23 27 13 16 18 15 12 20 17
## 2.3 2.31 2.32 2.33 2.34 2.35 2.36 2.37 2.38 2.39 2.4 2.41 2.42 2.43 2.44
## 21 13 16 9 5 7 8 6 8 7 13 5 8 6 4
## 2.45 2.46 2.47 2.48 2.49 2.5 2.51 2.52 2.53 2.54 2.55 2.56 2.57 2.58 2.59
## 4 3 3 9 3 17 17 9 8 9 3 3 3 3 1
## 2.6 2.61 2.63 2.64 2.65 2.66 2.67 2.68 2.7 2.71 2.72 2.74 2.75 2.77 2.8
## 3 3 3 1 1 3 1 2 1 1 3 3 2 1 2
## 3 3.01 3.02 3.04 3.05 3.11 3.22 3.24 3.4 3.5 3.51 3.65 3.67 4 4.01
## 8 14 1 2 1 1 1 1 1 1 1 1 1 1 2
## 4.13 4.5 5.01
## 1 1 1
```

### Problem 3: Data Wrangling Review

The assignment is to review Garrett Grolemond's "Data Wrangling with R" slides. This was done.

### Problem 4: Gapminder

The assignment here is to take a dataset from the Gapminder website ([www.gapminder.org](http://www.gapminder.org)) and create 2-5 plots making use of the histogram / frequency polygon / boxplot techniques taught in this lesson.

The dataset used is total cell phones per country per year. 275 countries are listed, with cell phone usage for each country from 1965 to 2011.

The question I wanted to find out from this data was, what regions of the world show the most cell phone usage and show the largest increase in cell phone usage?

The first step was to add another column of data for the 275 entries called "region." The following regions were assigned:

- 1 North America
- 2 South America / Central America
- 3 Europe
- 4 Asia
- 5 Middle East
- 6 Africa
- 7 Russia and former Soviet Republics ("Central Asia")
- 8 Australia / Oceania
- 9 Caribbean

The next step was to read in the data. So they wouldn't skew the results, any countries that showed no cell phone usage in 2011 were removed. Also, since the first cell phone usage numbers start to appear in 1980, only columns after 1980 are selected. Finally, column 1 was renamed to give a more meaningful name:

```
# Set up Libraries
library(dplyr)

##
## Attaching package: 'dplyr'
##
## The following objects are masked from 'package:stats':
##
##   filter, lag
##
## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union

library(ggplot2)

# Read data file
cp <- read.csv("cell_phone_total.csv", header = TRUE)

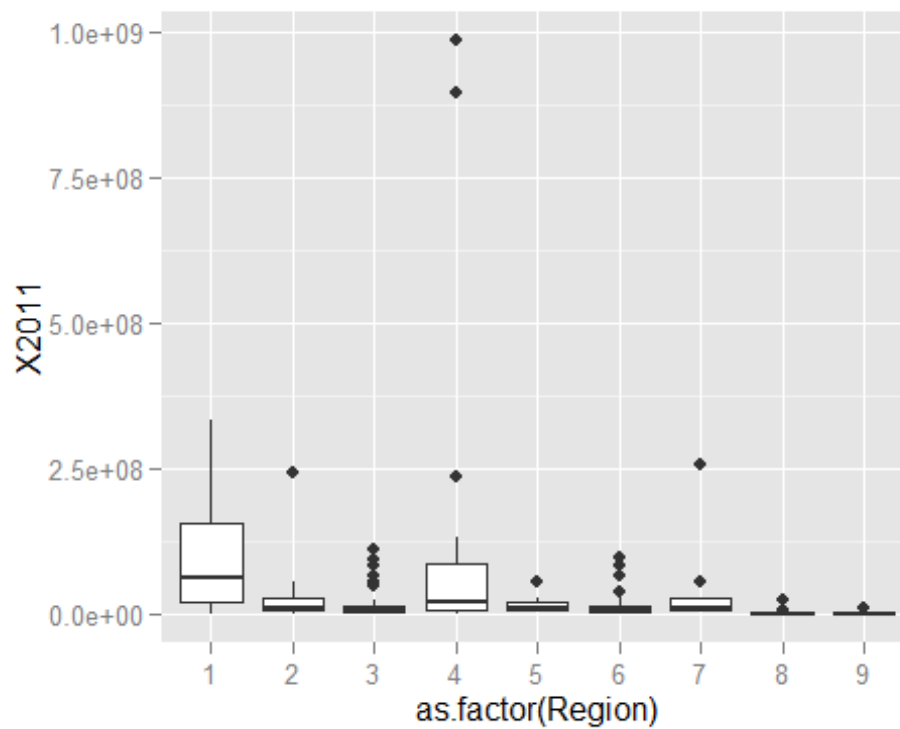
# Rename column 1 title
colnames(cp)[1] = "Country"

# Filter out any countries without 2011 data, any data before 1980
cpf <- cp %>% filter(!is.na(X2011))
cpf <- cpf %>% select(Country, Region,
                     starts_with("X198"), starts_with("X199"),
                     starts_with("X20") )
```

First we try a boxplot for each region of the total cell phone usage in the latest year (2011). To generate a valid boxplot, the "Region" column has to be converted from a numeric to a factor:


```
# Compare overall usage for most recent year for each region
qplot(x=as.factor(Region), y=X2011, data=cpf, geom="boxplot")
```





It is hard to interpret this data because of the very large scope of the data. We can get a slightly better sense by using log scales to get a sense of order-of-magnitude usage:


```
qplot(x=as.factor(Region), y=log10(X2011), data=cpf, geom="boxplot")
```



Here we get a slightly better sense of things. Looking at the medians, the region with the highest median use was North America (no surprise there), followed by Asia, with South/Central America, the Middle East, and (somewhat surprisingly) Russia with the same value, and Europe and Africa slightly behind them. Oceania and the Caribbean are significantly lower (also not surprising). Looking at the overall ranges, Asia has the country with the highest usage, followed by North America and Russia.

A lot of this might reflect the difference in populations of the countries. To eliminate this, another column showing the year-on-year increase from 2010 to 2011 ("pctIncrease") was added and a new boxplot was generated:

```
# Generate year-on-year increase for 2011 and do a boxplot by region
cpf <- mutate(cpf, pctIncrease = ((X2011 - X2010)/X2010)*100)
qplot(x=as.factor(Region), y=pctIncrease, data=cpf, geom="boxplot")
## Warning: Removed 1 rows containing non-finite values (stat_boxplot).
```




Here we see the median and middle quantiles grouped closer together in the 10-25% range, with Asia and Africa showing the largest median growth. The regions with the highest maximum values are Asia, Africa, and Australia.

We can also calculate the total for each region and do a simple plot:

```
# Calculate total usage by region
regions <- cpf %>% group_by(Region) %>%
  summarise(total=sum(as.numeric(X2011)))
regions

## Source: local data frame [9 x 2]
##
##   Region      total
##   (int)      (dbl)
## 1      1  452082365
## 2      2  516274151
## 3      3  688307902
## 4      4 2983362972
## 5      5  188950217
## 6      6  653127967
## 7      7  390841464
## 8      8   34023961
## 9      9   24076904

plot(regions)
```




The `as.numeric()` was used to prevent an overflow for the value of Asia. One can see from the graph and the table that Asia has by far the most overall users (with it containing China and India that makes sense), followed by Europe and Africa, South and Central America, North America, and Russia. The Middle East, Oceania and the Caribbean were a good deal behind. Two surprises were the relatively low ranking of North America (taking into account overall country populations that is a little less of a surprise), and the relatively high usage for Russia. Drilling down into the data, one can see that Russia itself has quite a few cell phone users (250m), with Ukraine contributing another 50m and Uzbekistan contributing another 25m. Given the population of Russia is around 143m, the 250m value is a particularly surprising number.

Finally, leaving out the regions, one can do a simple histogram to get a sense of the the distribution of cell phone usage among the countries:

```
qplot(x=log10(X2011), data=cpf)

## stat_bin: binwidth defaulted to range/30. Use 'binwidth = x' to adjust
this.
```



One can see that most of the countries fall in the  $10^7$  (10,000,000) range. The largest is just under 1 billion (China) and the smallest is in the 1000s (Tuvalu).

## Birthdays

The task is to investigate my friends' birthdays on Facebook. The goal is to find out 2 things: Which day has the most friends' birthdays, and what month has the most friends' birthdays?

Data was extracted by exporting the birthdays from Facebook to an Outlook calendar. The "export" feature in Outlook was then used to output the results to a CSV file. This was then read into R using "read.csv":

```
# Set up libraries
#install.packages("lubridate")
library(dplyr)
library(ggplot2)
library(lubridate)


## Warning: package 'lubridate' was built under R version 3.2.3

# Read data file
bd <- read.csv("friends_birthdays.csv", header = TRUE)
```

"lubridate" will be explained later.

To get a sense of which day has the most birthdays, first we do a simple plot of the date, where the date column is currently treated as a factor:

```
# Find birthday most commonly used
qplot(x = Start.Date, data=bd, binwidth = 1)
```



From this we see there is a 14-way tie for first, each having 3 friends with that birthday. To find out those days we can do a `group_by()` followed by a `summarise()` to get the count per day and then a `filter()` to only show those days with a value of 3:

```
bd %>% group_by(Start.Date) %>% summarise(num_friends = n()) %>%
  filter(num_friends == 3)
```

```
## Source: local data frame [14 x 2]
```

```
##
```

```
##   Start.Date num_friends
##   (fctr)      (int)
```

```
## 1 10/15/2016      3
```

```
## 2 11/21/2016      3
```

```
## 3  2/5/2016       3
```

```
## 4  3/1/2016       3
```

```
## 5  3/18/2016      3
```

```
## 6  3/25/2016      3
```

```
## 7  3/29/2016      3
```

```
## 8  4/16/2016      3
```

```
## 9  5/27/2016      3
```

```
## 10 6/20/2016      3
```

```
## 11 7/12/2016      3
## 12 7/20/2016      3
## 13 7/21/2016      3
## 14 8/27/2016      3
```

Finding the distribution by month is a little more involved. We use the "lubridate" package recommended by the Udacity team to process the date column. This library has already been installed and loaded from previous commands. We then use the "mdy()" function in lubridate to convert the "date" column that was formerly factors into official R dates (mdy means that the factor data is stored as month/day/year):

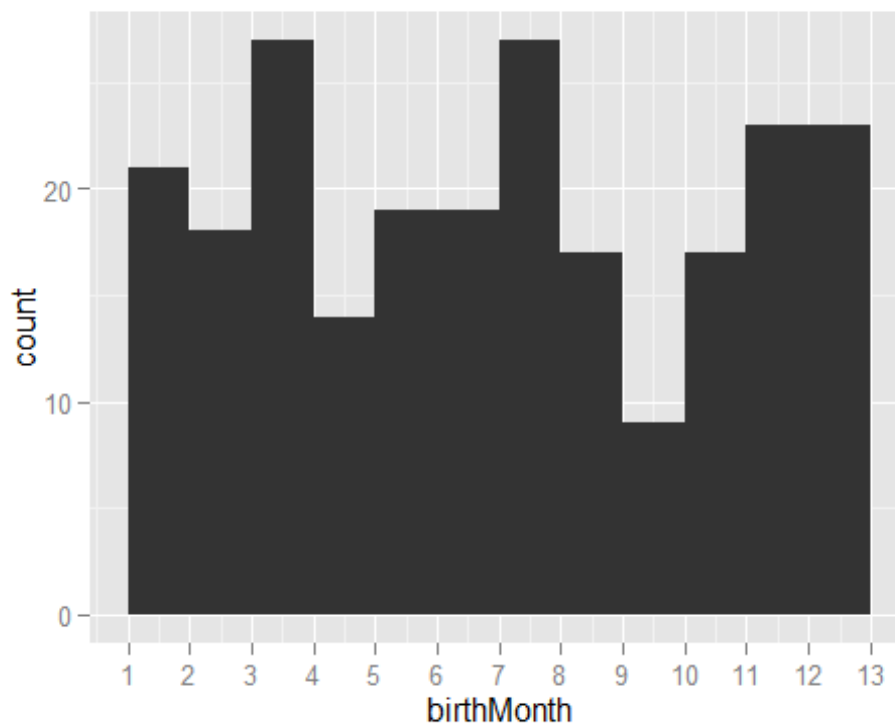
```
# Add month column
bd[,2] <- mdy(bd[,2])
```

Once we do this, we can then use the month() function to extract the month:

```
bd <- mutate(bd, birthMonth = month(Start.Date))
```

And we can then do a histogram by month:

```
qplot(x = birthMonth, data=bd, binwidth = 1) +
  scale_x_continuous(limit = c(1, 13), breaks = seq(1, 13, 1))
```



From this we see March and July have the most friends' birthdays, with 24 friends each.